

# **LAPORAN TUGAS AKHIR REKAYASA ORGANISASI DAN BIG DATA**



**Disusun oleh :  
Imam Akbar Nugraha - 1305213056**

**PROGRAM STUDI SARJANA SAINS DATA  
FAKULTAS INFORMATIKA  
UNIVERSITAS TELKOM  
BANDUNG  
2024**

## 1. Install transformers pinecone-client torch pandas

```
!pip install transformers pinecone-client torch pandas
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.3.0+cu121)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.0.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.14.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.23.3)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.25.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (24.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.1)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2024.5.15)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.31.0)
Requirement already satisfied: tokenizers<0.20,>=0.19 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.19.1)
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.4.3)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (4.66.4)
Requirement already satisfied: certifi>=2019.11.17 in /usr/local/lib/python3.10/dist-packages (from pinecone-client) (2024.6.2)
Collecting pinecone-plugin-interface<0.0.8,>=0.0.7 (from pinecone-client)
  Downloading pinecone_plugin_interface-0.0.7-py3-none-any.whl (6.2 kB)
Requirement already satisfied: typing-extensions>=3.7.4 in /usr/local/lib/python3.10/dist-packages (from pinecone-client) (4.12.2)
Requirement already satisfied: urllib3>=1.26.0 in /usr/local/lib/python3.10/dist-packages (from pinecone-client) (2.0.7)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch) (1.12.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch) (3.3)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch) (3.1.4)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch) (2023.6.0)
Collecting nvidia-cuda-nvrtc-cu12==12.1.105 (from torch)
  Using cached nvidia_cuda_nvrtc_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (23.7 MB)
Collecting nvidia-cuda-runtime-cu12==12.1.105 (from torch)
  Using cached nvidia_cuda_runtime_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (823 kB)
Collecting nvidia-cuda-cupti-cu12==12.1.105 (from torch)
  Using cached nvidia_cuda_cupti_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (14.1 MB)
Collecting nvidia-cudnn-cu12==8.9.2.26 (from torch)
  Using cached nvidia_cudnn_cu12-8.9.2.26-py3-none-manylinux1_x86_64.whl (731.7 MB)
Collecting nvidia-cublas-cu12==12.1.3.1 (from torch)
```

Baris perintah `!pip install transformers pinecone-client torch pandas` digunakan untuk menginstal paket transformers, pinecone-client, torch, dan pandas menggunakan pip. Paket transformers menyediakan berbagai model transformer seperti BERT, sementara pinecone-client adalah klien Python untuk layanan Pinecone yang digunakan untuk membuat dan mengelola indeks vektor. torch adalah pustaka yang mendukung komputasi numerik dan pembelajaran mesin, khususnya dalam pemrosesan model BERT. pandas digunakan untuk manipulasi dan analisis data dalam bentuk DataFrame.

## 2. Import Modul yang Diperlukan

```
import pandas as pd
import pinecone
from transformers import BertTokenizer, BertModel
import torch
import time
```

Pada bagian ini, modul yang diperlukan diimpor: pandas untuk manipulasi data, pinecone untuk berinteraksi dengan layanan Pinecone, transformers untuk menggunakan tokenizer dan model BERT, torch untuk manipulasi tensor dan komputasi model, serta time untuk mengukur waktu eksekusi. Modul-modul ini menyediakan fungsi dan kelas yang diperlukan untuk menjalankan tugas-tugas berikutnya dalam notebook.

### 3. Membaca Dataset

```
# Membaca data dari CSV
df = pd.read_csv('/content/spam.csv', sep=',', encoding='latin-1')
df
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN
...	...	...	...	...	...
5567	spam	This is the 2nd time we have tried 2 contact u...	NaN	NaN	NaN
5568	ham	Will l_b going to esplanade fr home?	NaN	NaN	NaN
5569	ham	Pity, * was in mood for that. So...any other s...	NaN	NaN	NaN
5570	ham	The guy did some bitching but I acted like i'd...	NaN	NaN	NaN
5571	ham	Rofl. Its true to its name	NaN	NaN	NaN

5572 rows × 5 columns

Kode `df = pd.read_csv('/content/spam.csv', sep=',', encoding='latin-1')` digunakan untuk membaca file CSV yang berisi data spam dan memuatnya ke dalam DataFrame `df`. Parameter `sep=','` menunjukkan bahwa pemisah kolom adalah koma, dan `encoding='latin-1'` digunakan untuk memastikan bahwa karakter dalam file CSV dapat dibaca dengan benar. DataFrame `df` kemudian digunakan untuk menyimpan dan memanipulasi data dalam bentuk tabel.

### 4. Inisialisasi Tokenizer dan Model BERT

```
# Inisialisasi tokenizer dan model BERT
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertModel.from_pretrained('bert-base-uncased')

# Fungsi untuk mendapatkan embedding teks
def get_text_embedding(text):
    inputs = tokenizer(text, return_tensors='pt', truncation=True, padding=True, max_length=512)
    with torch.no_grad():
        outputs = model(**inputs)
    embeddings = outputs.last_hidden_state.mean(dim=-1)
    return embeddings.detach().numpy().flatten().tolist()

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89: UserWarning:
The secret 'HF_TOKEN' does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 675B/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 3.97MB/s]
tokenizer.json: 100% 466k/466k [00:00<00:00, 6.00MB/s]
/usr/local/lib/python3.10/dist-packages/huggingface_hub/file_download.py:1132: FutureWarning: 'resume_download' is deprecated and will be removed in version 1.0.0. Downloads always resume when
warnings.warn(
config.json: 100% 570/570 [00:00<00:00, 5.77kB/s]
model.safetensors: 100% 440M/440M [00:03<00:00, 101MB/s]
```

Tokenizer dan model BERT diinisialisasi menggunakan model pra-terlatih `bert-base-uncased` dari pustaka `transformers`. Fungsi `get_text_embedding` digunakan untuk mendapatkan embedding teks dengan mengubah teks menjadi token menggunakan tokenizer, kemudian memproses token tersebut dengan model BERT

untuk menghasilkan tensor embedding. Keluaran dari BERT diambil dari `last_hidden_state` dan dirata-rata untuk mendapatkan representasi embedding.

## 5. Inisialisasi Pinecone

```
# Inisialisasi Pinecone
from pinecone import Pinecone, ServerlessSpec

api_key = "c61bfa4d-6c77-45d7-9cde-3b525cea05de"
pc = Pinecone(api_key=api_key)
```

Pada bagian ini, klien Pinecone diinisialisasi dengan menggunakan `api_key` yang diberikan. `ServerlessSpec` digunakan untuk mengatur spesifikasi serverless untuk index yang akan dibuat. Pinecone adalah layanan yang menyediakan indeks vektor yang dapat digunakan untuk pencarian kemiripan teks berbasis embedding.

## 6. Membuat Index Pinecone

```
# Membuat index Pinecone
index_name = "spam-messages-index"
if index_name not in pc.list_indexes():
    pc.create_index(
        name=index_name,
        dimension=768, # Replace with your model dimensions
        metric="cosine", # Replace with your model metric
        spec=ServerlessSpec(
            cloud="aws",
            region="us-east-1"
        )
    )
```

Kode ini mendefinisikan nama index sebagai `spam-messages-index` dan mengecek apakah index dengan nama tersebut sudah ada di Pinecone. Jika belum ada, index dibuat dengan dimensi 768 (sesuai dengan dimensi embedding dari model BERT) dan metrik cosine similarity. `ServerlessSpec` digunakan untuk mengatur spesifikasi serverless dengan cloud aws dan region us-east-1.

## 7. Menghubungkan ke Index

```
# Menghubungkan ke index
index = pc.Index(index_name)
```

Baris kode `index = pc.Index(index_name)` digunakan untuk menghubungkan ke index `spam-messages-index` yang telah dibuat atau sudah ada. Ini memungkinkan kita untuk melakukan operasi pencarian dan manipulasi data pada index tersebut.

## 8. Memasukkan Data ke Pinecone

```
# Memasukkan data ke Pinecone
for i, row in df.iterrows():
    text = row['v2'] # Asumsikan kolom teks adalah 'v2'
    embedding = get_text_embedding(text)
    index.upsert([(str(i), embedding)])
```

Dalam bagian ini, setiap baris dalam DataFrame df diiterasi. Teks diambil dari kolom v2, kemudian embedding teks diperoleh menggunakan fungsi `get_text_embedding`. Data embedding kemudian dimasukkan ke dalam index Pinecone menggunakan metode `upsert`, yang meng-update atau menyisipkan data baru ke dalam index.

## 9. Melakukan Query pada Index

```
# Melakukan query pada index
query_texts = [
    "Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005...",
    "WINNER!! As a valued network customer you have been selected to receive a £900 prize reward!",
    "I HAVE A DATE ON SUNDAY WITH WILL!!",
    "URGENT! You have won a 1 week FREE membership in our £100,000 Prize Jackpot!",
    "I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, I've cried enough today.",
    "Oh k...i'm watching here:)",
    "Fine if that's the way u feel. That's the way it's gotta b"
]
```

Daftar teks query didefinisikan untuk digunakan dalam pencarian pada index Pinecone. Teks-teks ini merupakan contoh pesan spam yang akan dicari kemiripannya dalam index.

## 10. Fungsi untuk Mengukur Waktu Eksekusi Query

```
# Fungsi untuk mengukur waktu eksekusi query
def measure_query_time(query_text):
    query_embedding = get_text_embedding(query_text)
    start_time = time.time()
    query_response = index.query(vector=query_embedding, top_k=10)
    end_time = time.time()
    duration = end_time - start_time
    return duration, query_response
```

Fungsi `measure_query_time` didefinisikan untuk mengukur waktu yang diperlukan untuk menjalankan query pada Pinecone. Fungsi ini mengubah teks query menjadi embedding, mencatat waktu sebelum dan sesudah query, serta menghitung durasi eksekusi. Fungsi ini mengembalikan durasi eksekusi dan respons query dari Pinecone.

## 11. Menjalankan Query dan Mencatat Waktu Eksekusi

```
# Menjalankan query dan mencatat waktu eksekusi
for query_text in query_texts:
    duration, query_response = measure_query_time(query_text)
    print(f"Query: {query_text}")
    print(f"Waktu eksekusi: {duration} detik")
    for match in query_response['matches']:
        print(f"ID: {match['id']}, Score: {match['score']}")
```

Query: Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005...

Waktu eksekusi: 0.729658842086792 detik

ID: 1162, Score: 0.89864713

ID: 2, Score: 0.89864713

ID: 4075, Score: 0.819228

ID: 2595, Score: 0.809554

ID: 877, Score: 0.809554

ID: 4241, Score: 0.806566954

ID: 4067, Score: 0.802371562

ID: 2611, Score: 0.794367313

ID: 2832, Score: 0.794095933

ID: 3003, Score: 0.794095933

Query: WINNER!! As a valued network customer you have been selected to receive a £900 prize reward!

Waktu eksekusi: 0.026117324829101562 detik

ID: 4034, Score: 0.885393679

ID: 5232, Score: 0.885393679

ID: 8, Score: 0.879962444

ID: 4279, Score: 0.878916323

ID: 1873, Score: 0.866761446

ID: 2495, Score: 0.848759711

ID: 3696, Score: 0.841698885

ID: 2112, Score: 0.840568244

ID: 93, Score: 0.839175761

ID: 4626, Score: 0.837771893

Query: I HAVE A DATE ON SUNDAY WITH WILL!!

Waktu eksekusi: 0.020369768142700195 detik

Query: I HAVE A DATE ON SUNDAY WITH WILL!!

Waktu eksekusi: 0.020369768142700195 detik

ID: 14, Score: 1.00100017

ID: 5155, Score: 0.78828609

ID: 259, Score: 0.780409098

ID: 1661, Score: 0.775901139

ID: 651, Score: 0.774294

ID: 4509, Score: 0.769461453

ID: 2937, Score: 0.769382179

ID: 2655, Score: 0.765554607

ID: 1878, Score: 0.762715399

ID: 3467, Score: 0.759862602

Query: URGENT! You have won a 1 week FREE membership in our £100,000 Prize Jackpot!

Waktu eksekusi: 0.02213001251220703 detik

ID: 93, Score: 0.844981

ID: 4626, Score: 0.842395842

ID: 187, Score: 0.840842664

ID: 1001, Score: 0.840842664

ID: 4040, Score: 0.838838458

ID: 4010, Score: 0.838838458

ID: 3966, Score: 0.83523792

ID: 4843, Score: 0.83523792

ID: 1873, Score: 0.832195938

ID: 4650, Score: 0.830685377

```

Query: I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, I've cried enough today.
Waktu eksekusi: 0.018512964248657227 detik
ID: 10, Score: 0.943577588
ID: 4183, Score: 0.823076427
ID: 99, Score: 0.803334355
ID: 2874, Score: 0.802716732
ID: 3665, Score: 0.799038768
ID: 3301, Score: 0.797587812
ID: 5538, Score: 0.793954194
ID: 846, Score: 0.791764677
ID: 5525, Score: 0.783107162
ID: 5246, Score: 0.780059159
Query: Oh k...i'm watching here:)
Waktu eksekusi: 0.018034934997558594 detik
ID: 16, Score: 1.00191426
ID: 399, Score: 0.845362425
ID: 4770, Score: 0.834993958
ID: 1218, Score: 0.829329491
ID: 2563, Score: 0.829329491
ID: 4386, Score: 0.82886517
ID: 3031, Score: 0.823122561
ID: 3009, Score: 0.818938792
ID: 2667, Score: 0.818586349
ID: 5321, Score: 0.818319201
Query: Fine if that's the way u feel. That's the way it's got a b
Waktu eksekusi: 0.016928672790527344 detik
ID: 18, Score: 0.90038836
ID: 4920, Score: 0.798796058
ID: 601, Score: 0.788626432
ID: 1830, Score: 0.784810364
ID: 1587, Score: 0.782213569
ID: 1784, Score: 0.782213569
ID: 4777, Score: 0.780885398
ID: 3189, Score: 0.779765904
ID: 1311, Score: 0.778784335

```

Query: Fine if that's the way u feel. That's the way it's got a b

Waktu eksekusi: 0.016928672790527344 detik

```

ID: 18, Score: 0.90038836
ID: 4920, Score: 0.798796058
ID: 601, Score: 0.788626432
ID: 1830, Score: 0.784810364
ID: 1587, Score: 0.782213569
ID: 1784, Score: 0.782213569
ID: 4777, Score: 0.780885398
ID: 3189, Score: 0.779765904
ID: 1311, Score: 0.778784335
ID: 1509, Score: 0.777925074

```

Pada bagian ini, query dijalankan untuk setiap teks dalam `query_texts`. Waktu eksekusi dicatat dan ditampilkan untuk setiap query. Hasil query, termasuk ID dan skor dari setiap pencocokan, juga ditampilkan. Hal ini memungkinkan untuk melihat performa dan hasil dari pencarian kemiripan teks pada index Pinecone.

## 12. Optimasi Query

```

# Optimasi Query
# Pinecone tidak memiliki parameter ef_search, tetapi Anda dapat menggunakan filter atau menyesuaikan jumlah hasil
optimized_query_texts = [
    "Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005...",
    "WINNER!! As a valued network customer you have been selected to receive a £900 prize reward!",
    "I HAVE A DATE ON SUNDAY WITH WILL!!",
    "URGENT! You have won a 1 week FREE membership in our £100,000 Prize Jackpot!",
    "I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, I've cried enough today.",
    "Oh k...i'm watching here:)",
    "Fine if that's the way u feel. That's the way it's got a b"
]

```



Daftar teks query yang dioptimasi didefinisikan. Optimasi ini bisa mencakup penggunaan filter atau parameter tambahan dalam query Pinecone untuk meningkatkan performa atau akurasi pencarian.

### 13. Menjalankan Query Optimasi dan Mencatat Waktu Eksekusi

```
# Menjalankan query optimasi dan mencatat waktu eksekusi
for query_text in optimized_query_texts:
    duration, query_response = measure_query_time(query_text)
    print(f"Query (Optimasi): {query_text}")
    print(f"Waktu eksekusi setelah optimasi: {duration} detik")
    for match in query_response['matches']:
        print(f"ID: {match['id']}, Score: {match['score']}")
```

Query (Optimasi): Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005...

Waktu eksekusi setelah optimasi: 0.1882789134979248 detik

ID: 1162, Score: 0.89864713

ID: 2, Score: 0.89864713

ID: 4075, Score: 0.819228

ID: 2595, Score: 0.809554

ID: 877, Score: 0.809554

ID: 4241, Score: 0.806566954

ID: 4067, Score: 0.802371562

ID: 2611, Score: 0.794367313

ID: 2832, Score: 0.794095933

ID: 3003, Score: 0.794095933

Query (Optimasi): WINNER!! As a valued network customer you have been selected to receive a £900 prize reward!

Waktu eksekusi setelah optimasi: 0.025378704071044922 detik

ID: 4034, Score: 0.885393679

ID: 5232, Score: 0.885393679

ID: 8, Score: 0.879962444

ID: 4279, Score: 0.878916323

ID: 1873, Score: 0.866761446

ID: 2495, Score: 0.848759711

ID: 3696, Score: 0.841698885

ID: 2112, Score: 0.840568244

ID: 93, Score: 0.839175761

ID: 4626, Score: 0.837771893

Query (Optimasi): I HAVE A DATE ON SUNDAY WITH WILL!!

Waktu eksekusi setelah optimasi: 0.019684314727783203 detik

ID: 14, Score: 1.00100017

ID: 5155, Score: 0.78828609

ID: 259, Score: 0.780409098

ID: 1661, Score: 0.775901139

ID: 651, Score: 0.774294

ID: 4509, Score: 0.769461453

ID: 2937, Score: 0.769382179

ID: 2655, Score: 0.765554607

ID: 1878, Score: 0.762715399

ID: 3467, Score: 0.759862602

Query (Optimasi): URGENT! You have won a 1 week FREE membership in our £100,000 Prize Jackpot!

Waktu eksekusi setelah optimasi: 0.019484996795654297 detik

ID: 93, Score: 0.844981

ID: 4626, Score: 0.842395842

ID: 187, Score: 0.840842664

ID: 1001, Score: 0.840842664

ID: 4040, Score: 0.838838458

ID: 4010, Score: 0.838838458

ID: 3966, Score: 0.83523792

ID: 4843, Score: 0.83523792

ID: 1873, Score: 0.832195938

ID: 4650, Score: 0.830685377



```
Query (Optimasi): I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, I've cried enough today.
Waktu eksekusi setelah optimasi: 0.022426605224609375 detik
ID: 10, Score: 0.943577588
ID: 4183, Score: 0.823076427
ID: 99, Score: 0.803334355
ID: 2874, Score: 0.802716732
ID: 3665, Score: 0.799038768
ID: 3301, Score: 0.797587812
ID: 5538, Score: 0.793954194
ID: 846, Score: 0.791764677
ID: 5525, Score: 0.783107162
ID: 5246, Score: 0.780059159
Query (Optimasi): Oh k...i'm watching here:)
Waktu eksekusi setelah optimasi: 0.020049333572387695 detik
ID: 16, Score: 1.00191426
ID: 399, Score: 0.845362425
ID: 4770, Score: 0.834993958
ID: 1218, Score: 0.829329491
ID: 2563, Score: 0.829329491
ID: 4386, Score: 0.82886517
ID: 3031, Score: 0.823122561
ID: 3009, Score: 0.818938792
ID: 2667, Score: 0.818586349
ID: 5321, Score: 0.818319201
Query (Optimasi): Fine if that's the way u feel. That's the way it's got a b
Waktu eksekusi setelah optimasi: 0.023157596588134766 detik
ID: 18, Score: 0.90038836
ID: 4920, Score: 0.798796058
ID: 601, Score: 0.788626432
ID: 1830, Score: 0.784810364
ID: 1587, Score: 0.782213569
ID: 1784, Score: 0.782213569
ID: 4777, Score: 0.780885398
ID: 3189, Score: 0.779765904
ID: 1311, Score: 0.778784335
```

```
Query (Optimasi): Fine if that's the way u feel. That's the way it's got a b
Waktu eksekusi setelah optimasi: 0.023157596588134766 detik
ID: 18, Score: 0.90038836
ID: 4920, Score: 0.798796058
ID: 601, Score: 0.788626432
ID: 1830, Score: 0.784810364
ID: 1587, Score: 0.782213569
ID: 1784, Score: 0.782213569
ID: 4777, Score: 0.780885398
ID: 3189, Score: 0.779765904
ID: 1311, Score: 0.778784335
ID: 1509, Score: 0.777925074
```

Query dijalankan pada teks yang dioptimasi, dan waktu eksekusi dicatat serta ditampilkan untuk setiap query yang dioptimasi. Hasil query yang dioptimasi, termasuk ID dan skor dari setiap pencocokan, juga ditampilkan. Hal ini memungkinkan untuk membandingkan performa dan hasil antara query standar dan query yang dioptimasi.

# 14. Output Membuat Index pada Database Pinecone

Default

DATABASE

Indexes (1)

Collections

SOLUTIONS

Assistant (Beta)

MANAGE

API Keys

Members

STARTER USAGE

Storage

0 / 2GB

RUs

123 / 1M

WUs

131K / 2M

Upgrade now

Back to indexes

Connect

spam-messages-index

METRIC

cosine

DIMENSIONS

768

HOST

https://spam-messages-index-7vtkmz1.svc.aped-4627-b74a.pinecone.io

CLOUD

AWS

REGION

us-east-1

TYPE

Serverless

VECTOR COUNT

5,572

BROWSER

METRICS

NAMESPACES (1)

Namespace

( Default )

Query by

dense vector value

vector

0.48,0.3,0.4,0.05,0.71,0.56,0.71,0.86,0.73,0.28,0.71,0.45,0.

Top K

10

Query

Metadata Filter

Default

DATABASE

Indexes (1)

Collections

SOLUTIONS

Assistant (Beta)

MANAGE

API Keys

Members

STARTER USAGE

Storage

0 / 2GB

RUs

123 / 1M

WUs

131K / 2M

Upgrade now

Matches: 10

+ Upsert Record

1

ID

3604

VALUES

-0.0842527673, -0.825902283, 0.0180563927, 0.181997091, -0.293622375, -0.0606717095, ...

SCORE

0.0227

METADATA

2

ID

723

VALUES

0.109117486, 0.530179858, 0.311076939, 0.287302792, -0.0565993898, -0.15794988, 0.443...

SCORE

0.0166

METADATA

3

ID

5383

VALUES

0.161717966, 0.27261728, 0.631023943, -0.417316735, 0.154481396, 0.00626505492, 0.742...

SCORE

0.0164

METADATA

4

ID

3934

VALUES

0.0715865493, -0.0365648344, 0.16315715, -0.179758251, 0.0268870592, -0.292345464, 0....

SCORE

0.0164

METADATA

Default

DATABASE

Indexes (1)

Collections

SOLUTIONS

Assistant (Beta)

MANAGE

API Keys

Members

STARTER USAGE

Storage

0 / 2GB

RUs

123 / 1M

WUs

131K / 2M

Upgrade now

4

ID

3934

VALUES

0.0715865493, -0.0365648344, 0.16315715, -0.179758251, 0.0268870592, -0.292345464, 0....

SCORE

0.0164

METADATA

5

ID

309

VALUES

0.254617691, 0.463602871, 0.162182406, 0.272158623, 0.822626889, -0.493521929, -0.214...

SCORE

0.0162

METADATA

6

ID

1700

VALUES

0.409445047, -0.0643135086, 0.557512403, -0.340661, -0.0420853645, -0.11097309, 0.201...

SCORE

0.0153

METADATA

7

ID

2789

VALUES

0.281101644, -0.0743034557, 0.512543857, -0.301090628, 0.299027711, -0.128831059, 0.58...

SCORE

0.0147

METADATA

ID	SCORE	VALUES
7 2789	0.0147	0.281101644, -0.0743034557, 0.512543857, -0.301090628, 0.299027711, -0.128831059, 0.58...
8 2087	0.0142	0.619081199, 0.107771777, 0.49450174, -0.0593207702, 0.184056014, -0.12985079, 0.27422...
9 4366	0.0142	-0.0230318867, 0.0867015123, 0.646155119, -0.0142318644, 0.0926808864, -0.315524697, ...
10 959	0.0132	0.0878898352, 0.208830699, -0.324097365, -0.183845982, 0.110290252, -0.282553256, 0...

Query used 6 RUs

Tampilan dashboard Pinecone digunakan untuk mengelola indeks vektor, dengan detail indeks bernama `spam-messages-index`, yang dirancang untuk mengelola data terkait pembelian pulsa dan pengecekan kuota. Metrik yang digunakan adalah cosine untuk mengukur kesamaan antar vektor, dan dimensi indeks adalah 768, sesuai dengan panjang vektor embedding dari model seperti BERT. Indeks ini di-host di AWS, dalam region `us-east-1`, menggunakan tipe serverless untuk kemudahan skalabilitas dan manajemen infrastruktur. Pengguna dapat memasukkan vektor query langsung di antarmuka, menentukan jumlah hasil teratas (Top K) yang diinginkan, dan menjalankan pencarian dengan menekan tombol Query. Hasil pencarian ditampilkan dengan skor relevansi dan ID dari setiap hasil, menunjukkan tingkat kesamaan dengan vektor query yang dimasukkan. Setiap hasil memiliki ID unik dan skor relevansi yang menunjukkan kemiripannya dengan query, serta menampilkan sebagian dari nilai vektor terkait hasil pencarian. Tampilan ini memperlihatkan bagaimana Pinecone memudahkan manajemen dan pencarian dalam indeks vektor, memungkinkan pengguna untuk menjalankan pencarian berbasis kesamaan vektor dengan mudah dan mendapatkan hasil relevan berdasarkan metrik yang telah ditentukan. Sistem ini sangat berguna untuk berbagai aplikasi seperti pencarian teks, rekomendasi produk, dan analisis data berbasis vektor lainnya.

**Link Google Collabs :**

<https://colab.research.google.com/drive/1z4TBI426q99Qvx4nN014rb1lcWjZnhrB?usp=sharing>

**Link Video Presentasi :**

[https://drive.google.com/drive/folders/1cJ8b\\_dXraru5BEWNyK\\_WDDzoFif5\\_A2O?usp=sharing](https://drive.google.com/drive/folders/1cJ8b_dXraru5BEWNyK_WDDzoFif5_A2O?usp=sharing)